

class 8 halloween

```
candy_file <- read.csv("candy-data.csv", row.names = 1)

head(candy_file)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
nrow(candy_file)
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset? The functions `dim()`, `nrow()`, `table()` and `sum()` may be useful for answering the first 2 questions.

```
dim(candy_file)
```

```
[1] 85 12
```

```
sum(candy_file$fruity)
```

```
[1] 38
```

Q3. What is your favorite candy in the dataset and what is it's winpercent value?

```
candy_file["Almond Joy",]$winpercent
```

```
[1] 50.34755
```

Q4. What is the winpercent value for “Kit Kat”?

```
candy_file["Kit Kat",]$winpercent
```

```
[1] 76.7686
```

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy_file["Tootsie Roll Snack Bars",]$winpercent
```

```
[1] 49.6535
```

```
library(skimr)
skim(candy_file)
```

Table 1: Data summary

		candy_file
Name		
Number of rows		85
Number of columns		12
Column type frequency:		
numeric		12

Group variables

None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

```
skimr::skim(candy_file)
```

Table 3: Data summary

Name	candy_file
Number of rows	85
Number of columns	12
Column type frequency: numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

hist

Q7. What do you think a zero and one represent for the candy\$chocolate column? 0 - no chocolate 1- chocolate

```
candy_file$chocolate
```

```
[1] 1 1 0 0 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 1 1 1 1 0 1 1 0 0 0 1 1 0 1 1 1
[39] 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 1 1
[77] 1 1 0 1 0 0 0 0 1
```

```
choc.ind <- as.logical(candy_file$chocolate)
fruit.ind <- as.logical(candy_file$fruity)
choc.win <- candy_file[choc.ind,]$winpercent
choc.win
```

```
[1] 66.97173 67.60294 50.34755 56.91455 38.97504 55.37545 62.28448 56.49050
[9] 59.23612 57.21925 76.76860 71.46505 66.57458 55.06407 73.09956 60.80070
[17] 64.35334 47.82975 54.52645 70.73564 66.47068 69.48379 81.86626 84.18029
[25] 73.43499 72.88790 65.71629 34.72200 37.88719 76.67378 59.52925 48.98265
[33] 43.06890 45.73675 49.65350 81.64291 49.52411
```

```
mean(choc.ind)
```

```
[1] 0.4352941
```

```
mean(fruit.ind)
```

```
[1] 0.4470588
```

```
mean(candy_file$chocolate)
```

```
[1] 0.4352941
```

```
mean(candy_file$fruity)
```

```
[1] 0.4470588
```

Q8. Plot a histogram of winpercent values

```
hist(candy_file$winpercent)
```



Q9. Is the distribution of winpercent values symmetrical? NO

Q10. Is the center of the distribution above or below 50%?

```
candy_file$winpercent[as.logical(candy_file$nougat)]
```

```
[1] 67.60294 56.91455 38.97504 73.09956 60.80070 46.29660 76.67378
```

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
mean(candy_file$chocolate) >= mean(candy_file$fruity)
```

```
[1] FALSE
```

Q12. Is this difference statistically significant?

Welch Two Sample t-test

data: chocolate and fruity

t = 6.2582, df = 68.882, p-value = 2.871e-08

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

11.44563 22.15795

sample estimates:

mean of x mean of y

60.92153 44.11974

Q13. What are the five least liked candy types in this set?

```
x <- c(5,2,3,6)
sort(x)
```

```
[1] 2 3 5 6
```

```
sort(x, decreasing=TRUE)
```

```
[1] 6 5 3 2
```

```
sort(x, decreasing =FALSE)
```

```
[1] 2 3 5 6
```

```
x
```

```
[1] 5 2 3 6
```

```
order(x)
```

```
[1] 2 3 1 4
```

```
x[order(x)]
```

```
[1] 2 3 5 6
```

```
y<-c("D", "A", "E")  
order(y)
```

```
[1] 2 1 3
```

Q14. What are the top 5 all time favorite candy types out of this set?

```
order.win <- order(candy_file$winpercent)  
candy_file[order.win[1:5],1]
```

```
[1] 0 0 0 0 0
```

```
tail(order.win)
```

```
[1] 54 65 29 80 52 53
```

```
head(order.win)
```

```
[1] 45 8 13 73 27 58
```

```
head(candy_file[order(candy_file$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0

	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip				0	0	0	1	0.197		0.976
Boston Baked Beans				0	0	0	1	0.313		0.511
Chiclets				0	0	0	1	0.046		0.325
Super Bubble				0	0	0	0	0.162		0.116
Jawbusters				0	1	0	1	0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

```
tail(candy_file[order(candy_file$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0

	crisped	rice	wafer	hard	bar	pluribus	sugarpercent
Snickers		0	0	1		0	0.546
Kit Kat		1	0	1		0	0.313
Twix		1	0	1		0	0.546
Reese's Miniatures		0	0	0		0	0.034
Reese's Peanut Butter cup		0	0	0		0	0.720


	pricepercent	winpercent
Snickers	0.651	76.67378
Kit Kat	0.511	76.76860
Twix	0.906	81.64291
Reese's Miniatures	0.279	81.86626
Reese's Peanut Butter cup	0.651	84.18029

```
x <- order(candy_file$winpercent)
order(candy_file$winpercent)
```

```
[1] 45  8 13 73 27 58 72  3 71 20 10 70 60 56 12 51 49 63  9 11 82 31 17 46 15
[26] 50 30 84 22 14 59 76 16 83 81 77 64  4 47 35 18 79 40 75 85 78  6 21  5 68
[51] 32 41 74 36 62 42 23 25  7 19 28 26 66 67 38 24 61 39 57 44 34  1 69  2 48
[76] 43 33 55 37 54 65 29 80 52 53
```

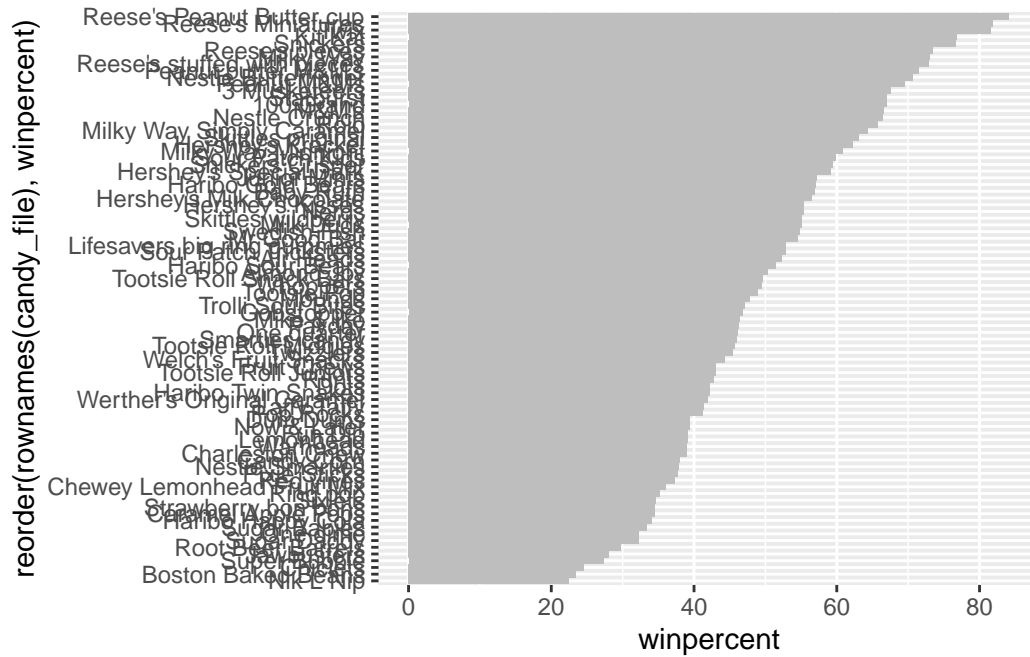
Q15. Make a first barplot of candy ranking based on winpercent values.

```
library(ggplot2)
ggplot(candy_file) +
  aes(winpercent, rownames(candy_file))+
  geom_col()
```



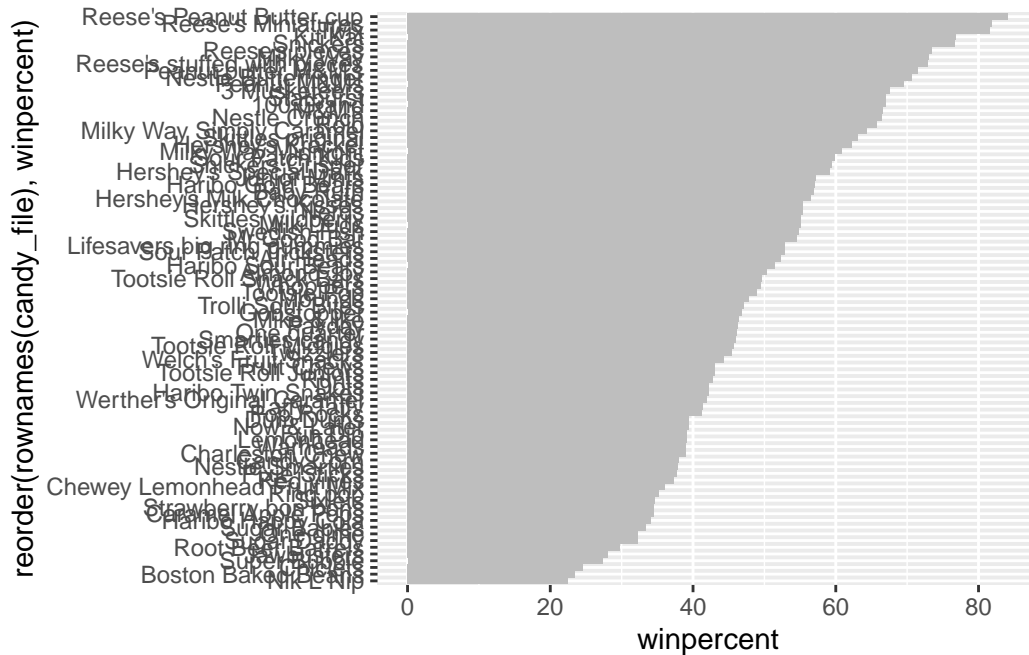
Q16. This is quite ugly, use the `reorder()` function to get the bars sorted by winpercent?

```
ggplot(candy_file) +
  aes(winpercent, rownames(candy_file), reorder(candy_file, FUN=mean(candy_file$winpercent)))
  geom_col()
```

```
mycols <- rep("gray",nrow(candy_file))

ggplot(candy_file) +
  aes(winpercent, reorder(rownames(candy_file),winpercent)) +
  geom_col(bg=mycols)
```




```
mycols <- rep("gray", nrow(candy_file))
mycols[as.logical(candy_file$chocolate)] <- "brown"
mycols=rep("black", nrow(candy_file))
mycols[as.logical(candy_file$chocolate)] = "chocolate"
mycols[as.logical(candy_file$bar)] = "orange"
mycols[as.logical(candy_file$fruity)] = "pink"
mycols
```

```
[1] "orange" "orange" "black" "black" "pink" "orange"
[7] "orange" "black" "black" "pink" "orange" "pink"
[13] "pink" "pink" "pink" "pink" "pink" "pink"
[19] "pink" "black" "pink" "pink" "chocolate" "orange"
[25] "orange" "orange" "pink" "chocolate" "orange" "pink"
[31] "pink" "pink" "chocolate" "chocolate" "pink" "chocolate"
[37] "orange" "orange" "orange" "orange" "orange" "pink"
[43] "orange" "orange" "pink" "pink" "orange" "chocolate"
[49] "black" "pink" "pink" "chocolate" "chocolate" "chocolate"
[55] "chocolate" "pink" "chocolate" "black" "pink" "chocolate"
[61] "pink" "pink" "chocolate" "pink" "orange" "orange"
[67] "pink" "pink" "pink" "pink" "black" "black"
[73] "pink" "pink" "pink" "chocolate" "chocolate" "orange"
[79] "pink" "orange" "pink" "pink" "pink" "black"
```

```
[85] "chocolate"
```

```
ggplot(candy_file) +  
  aes(winpercent, reorder(rownames(candy_file), winpercent)) +  
  geom_col(bg=mycols)
```



Q17. What is the worst ranked chocolate candy? Resse's Peanut Butter cup

Q18. What is the best ranked fruity candy? Nik L Nip

How about a plot of price vs win


Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

```
library(ggrepel)

ggplot(candy_file) +
  aes(winpercent, pricepercent, label=rownames(candy_file)) +
  geom_point(col=mycols) +
```

```
geom_text_repel(col=mycols, size=3.3, max.overlaps = 5)
```

Warning: ggrepel: 65 unlabeled data points (too many overlaps). Consider increasing max.overlaps




Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

```
ord <- order(candy_file$pricepercent, decreasing = TRUE)
head( candy_file[ord,c(11,12)], n=5 )
```

	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Q21. Make a barplot again with `geom_col()` this time using `pricepercent` and then improve this step by step, first ordering the x-axis by value and finally making a so called “dot chat” or “lollipop” chart by swapping `geom_col()` for `geom_point()` + `geom_segment()`.

```
ggplot(candy_file) +
  aes(pricepercent, reorder(rownames(candy_file), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy_file), pricepercent),
                    xend = 0), col="gray40") +
  geom_point()
```




Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy_file)
corrplot(cij)
```

chocolate and fruity

Q23. Similarly, what two variables are most positively correlated? chocolate and bar abd
oricepercent and winpercent

#Principal Component Analysis


```
pca <- prcomp(candy_file, scale=TRUE)
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369


	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

```
plot(pca$x[,1:2], col=mycols, pch=16)
```



```
my_data <- cbind(candy_file, pca$x[,1:3])
p <- ggplot(my_data) +
  aes(x=PC1, y=PC2,
      size=winpercent/100,
      text=rownames(my_data),
      label=rownames(my_data)) +
  geom_point(col=mycols)
```

p




```
library(ggrepel)

p + geom_text_repel(size=3.3, col=mycols, max.overlaps = 7) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
        subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown)",
        caption="Data from 538")
```

Warning: ggrepel: 59 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),



Data from 538